

Frost64 CPU

Andrew Clark

June 27, 2020

1 Table of Contents

Contents

1	Table of Contents	1
2	Registers, Main Widths, etc.	2
3	Encoding Functional Unit Numbers Into Arithmetic/Logic Instructions	4
4	Instructions (CPU's perspective)	5
4.1	Group 0 Instructions	5
4.2	Group 1 Instructions	5
4.3	Group 2 Instructions	6
4.4	Group 3 Instructions	7
4.5	Group 4 Instructions	8
4.6	Group 5 Instructions	8
4.7	Group 6 Instructions	9
4.8	Group 7 Instructions	13

2 Registers, Main Widths, etc.

- The main width of the processor, hinted at with the name, is 64-bit. Addresses are 64-bit, and some instructions only support 64-bit operations.
- Instructions LARs (ILARs)
 - There are 256 total ILARs, with 128 of them used for interrupt service routines and 128 of them used for user code.
 - The two ILARs named `i0`, which are ILAR number 0 and ILAR number 128 within the ILAR file, have their fields always set to zero. They cannot be changed.
 - ILARs are 256 bytes long, and instructions are 32-bit. This means that one ILAR holds $256 / 4 = 64$ instructions.
- Data LARs (DLARs)
 - There are 256 total DLARs, with 128 used for interrupt service routines and 128 of them used for user code.
 - The two DLARs named `dzero`, which are DLAR number 0 and DLAR number 128 within the DLAR file, have their fields always set to zero. They cannot be changed.
 - The user-mode DLAR named `dira`, which is DLAR 1 within the DLAR file, is treated specially, as its address field is the interrupt return address for the `reti` instruction.
 - DLARs are 256 bytes long.
 - DLARs can take on the following type tags (3-bit enum):
 - * 8-bit, unsigned
 - * 8-bit, signed
 - * 16-bit, unsigned
 - * 16-bit, signed
 - * 32-bit, unsigned
 - * 32-bit, signed
 - * 64-bit, unsigned
 - * 64-bit, signed
 - The base address of a DLAR is $64 - 8 = 56$ bits long.
 - The scalar offset of a DLAR is 8 bits long.
- Program Counter (PC)
 - The program counter consists of an ILAR number (7-bit) and an offset into said ILAR (6-bit).
 - Two program counters exist: one for when servicing an interrupt, and one for when not servicing an interrupt (user mode).

- **Interrupt Enable (ie)**
 - This 1-bit register indicates whether or not an interrupt can be serviced.

3 Encoding Functional Unit Numbers Into Arithmetic/Logic Instructions

Here is an explanation of what encoding functional unit numbers into arithmetic/logic instructions means to the CPU.

”This instruction, which takes multiple cycles to complete, uses this functional unit. You may execute later instructions that don’t depend on the results that this functional unit will produce at the same time as this instruction. If a later instruction needs the results of this functional unit, please stall, as necessary, until the functional unit has produced its results.”

In other words, said instructions start an asynchronous use of the functional unit, stalling only if necessary, i.e. when a needed result isn’t computed yet.

Different implementations of the CPU may have a different number of functional units, with the machine still executing instructions correctly, even if it has fewer functional units than what was encoded into the instruction. Four bits in relevant instructions are used to indicate which functional unit to use for said instructions.

4 Instructions (CPU's perspective)

4.1 Group 0 Instructions

- Encoding: 000a aaaa aabb bbbb bccc cccc vooo ffff
 - a: DLAR a
 - b: DLAR b
 - c: DLAR c
 - v:
 - * when 0: scalar operation
 - * when 1: vector operation
 - o: opcode
 - f: functional unit number f
- Instruction List:
 1. add fF, dA, dB, dC
 2. sub fF, dA, dB, dC
 3. slt fF, dA, dB, dC
 4. mul fF, dA, dB, dC
 5. div fF, dA, dB, dC
 - Note: Perform an unsigned divide if dB is unsigned, but perform a signed divide if dB is signed.
 6. and fF, dA, dB, dC
 7. orr fF, dA, dB, dC
 8. xor fF, dA, dB, dC

4.2 Group 1 Instructions

- Encoding: 001a aaaa aabb bbbb bccc cccc vooo ffff
 - a: DLAR a
 - b: DLAR b
 - c: DLAR c
 - o: opcode
 - v:
 - * when 0: scalar operation
 - * when 1: vector operation
 - f: functional unit number f

- Instruction List:

1. `shl fF, dA, dB, dC`
 - Shift left
2. `shr fF, dA, dB, dC`
 - Shift right
 - Note: Perform a logical right shift if `dB` is unsigned, but perform an arithmetic right shift if `dB` is signed.
3. Reserved for future expansion.
4. Reserved for future expansion.
5. Reserved for future expansion.
6. Reserved for future expansion.
7. Reserved for future expansion.
8. Reserved for future expansion.

4.3 Group 2 Instructions

- Encoding: `010a aaaa aabb bbbb biii iiii iiii oooo`

- `a`: DLAR `a`
- `b`: DLAR `b`
- `i`: Signed 11-bit immediate
- `o`: opcode

- Notes:

- These instructions compute the address to load from or store to via the following formula: `u64(dB.scalar_data) + s64(simm11)`

- Instruction List:

1. `ldiu8 dA, [dB, simm11]`
2. `ldis8 dA, [dB, simm11]`
3. `ldiu16 dA, [dB, simm11]`
4. `ldis16 dA, [dB, simm11]`
5. `ldiu32 dA, [dB, simm11]`
6. `ldis32 dA, [dB, simm11]`
7. `ldiu64 dA, [dB, simm11]`
8. `ldis64 dA, [dB, simm11]`
9. `stiu8 dA, [dB, simm11]`
10. `stis8 dA, [dB, simm11]`

```

11. stiu16 dA, [dB, simm11]
12. stis16 dA, [dB, simm11]
13. stiu32 dA, [dB, simm11]
14. stis32 dA, [dB, simm11]
15. stiu64 dA, [dB, simm11]
16. stis64 dA, [dB, simm11]

```

4.4 Group 3 Instructions

- Encoding: 011a aaaa aabb bbbb bccc cccc iiii oooo
 - a: DLAR a
 - b: DLAR b
 - b: DLAR c
 - i: Signed 4-bit immediate
 - o: opcode
- Notes:
 - These instructions compute the address to load from or store to via the following formula: $u64(dB.scalar_data) + dC.addr + s64(simm4)$
- Instruction List:


```

1. ldu8 dA, [dB, dC, simm4]
2. lds8 dA, [dB, dC, simm4]
3. ldu16 dA, [dB, dC, simm4]
4. lds16 dA, [dB, dC, simm4]
5. ldu32 dA, [dB, dC, simm4]
6. lds32 dA, [dB, dC, simm4]
7. ldu64 dA, [dB, dC, simm4]
8. lds64 dA, [dB, dC, simm4]
9. stu8 dA, [dB, dC, simm4]
10. sts8 dA, [dB, dC, simm4]
11. stu16 dA, [dB, dC, simm4]
12. sts16 dA, [dB, dC, simm4]
13. stu32 dA, [dB, dC, simm4]
14. sts32 dA, [dB, dC, simm4]
15. stu64 dA, [dB, dC, simm4]
16. sts64 dA, [dB, dC, simm4]

```


4.5 Group 4 Instructions

- Encoding: 100a aaaa aabb bbbb bccc cccc iiii jjjj
 - a: ILAR a (base ILAR)
 - b: DLAR b
 - c: ILAR c
 - i: Signed 4-bit immediate (i_simm4)
 - j: Unsigned 4-bit immediate (j_imm4). How many additional, consecutive ILARs, starting with iA , will be fetched into.
 - * Note: The destination ILAR number will wrap around to $i0$ if incrementing from $i128$ to the next ILAR, i.e. a 7-bit counter is used.
- Notes:
 - The one instruction in this group of instructions, `fetch`, computes the starting address, via the following equation: $u64(dB.scalar_data) + iC.addr + s64(i_simm4)$
- Instruction List:
 - `fetch iA, dB, iC, i_simm4, j_imm4`

4.6 Group 5 Instructions

- Encoding: 101a aaaa aabb bbbb biii iiii iiij jjjo
 - a: ILAR a (base ILAR)
 - b: DLAR b or ILAR c
 - i: Signed 10-bit immediate (i_simm10)
 - j: Unsigned 4-bit immediate (j_imm4). How many additional, consecutive ILARs, starting with iA , will be fetched into.
 - * Note: The destination ILAR number will wrap around to $i0$ if incrementing from $i128$ to the next ILAR, i.e. a 7-bit counter is used.
 - o: Opcode
- Instruction List:
 1. `fetchd iA, dB, i_simm10, j_imm4`
 - This instruction acts almost the same as a `fetch` instruction, but computes the starting address to fetch from via the following formula:
 - * $u64(dB.scalar_data) + s64(i_simm10)$
 2. `fetcha iA, iB, i_simm10, j_imm4`
 - This instruction acts almost the same as a `fetch` instruction, but computes the starting address to fetch from via the following formula:
 - * $iB.addr + s64(i_simm10)$

4.7 Group 6 Instructions

- Encoding: 110a aaaa aabb bbbb biii iiio oooo ffff
 - a: DLAR a (mode-specific)
 - b: LAR b (instruction-specific, either a DLAR or an ILAR)
 - i: One of the following, depending on the instruction:
 - * 6-bit immediate, `imm6`
 - * 6-bit signed immediate, `simm6`
 - o: opcode
 - f: Functional unit number
- Instruction List:
 1. `gudata dA, dB, imm6`
 - Name: "get user-mode data"
 - For this instruction, the destination, mode-specific DLARs are treated as vectors of 128-bit data, where each vector element stores the following:
 - * The full address (64-bit) of user-mode DLARs (starting with `dB`), stored in vector element bit range `63 . . 0`.
 - * The type tag (3-bit) of user-mode DLARs (starting with `dB`), stored in vector element bit range `66 . . 64`
 - Each 128-bit vector element corresponds to one user-mode DLAR's contents.
 - This instruction grabs the full address and type tag of user-mode DLARs, starting with `dB`, and including the following (consecutive) `imm6` user-mode DLARs, storing the 128-bit vector elements in (mode-specific) DLARs, starting with `dA` as the base DLAR. The base mode-specific DLAR is followed by consecutive DLARs if needed to store the user-mode DLAR fields.
 - If the last mode-specific DLAR is not completely filled with user-mode DLAR contents, its remaining 128-bit vector elements will be filled with `0x0`.
 2. `rudata dA, dB, imm6`
 - Name: "restore user-mode data"
 - This instruction sort of does the reverse of the `gudata` instruction, performing the equivalent of load instructions, loading into user-mode DLARs based upon the metadata stored in mode-specific DLARs. As such, this instruction can read from memory, just like load instructions.
 3. `guinst dA, iB, imm6`
 - Name: "get user-mode instructions"

- For this instruction, the destination, mode-specific DLARs are treated as vectors of 64-bit data, where each vector element stores the address field of a single user-mode ILAR.
 - Each 64-bit vector element corresponds to one user-mode ILAR's address.
 - This instruction grabs the full address of user-mode ILARs, starting with `dB` and including the following (consecutive) `imm6` user-mode ILARs, storing the 64-bit vector elements in mode-specific DLARs, starting with `dA` as the base DLAR. The base mode-specific DLAR is followed by consecutive DLARs if needed to store the user-mode ILAR addresses.
 - If the last mode-specific DLAR is not completely filled with user-mode ILAR addresses, its remaining 64-bit vector elements will be filled with `0x0`.
4. `ruinst dA, iB, imm6`
- Name: "restore user-mode instructions"
 - This instruction performs `imm6 + 1` consecutive `fetch` operations, starting with (mode-specific) `dA` as the base register that is the source of addresses to `fetch` instructions from. User-mode `iB` is the first ILAR that instructions will be `fetch`d into.
 - Basically, this instruction takes the results from a `guinst` instruction and re-`fetch`s the instructions back into user-mode ILARs.
5. `gupc dA`
- Name: "get user-mode pc"
 - This instruction grabs the user-mode program counter value, encoded into a scalar value as
`* concat(ilar_number, ilar_scalar_offset),`
and stores that value into DLAR `dA`.
 - Note: the type tag of mode-specific DLAR `dA` is respected during this operation.
 - Note: This instruction is useful for implementing user-mode jump tables that are (at least partially) stored in a bunch of ILARs.
6. `rupc dA`
- Name: "restore user-mode pc"
 - This instruction simply does the reverse operation of the `gupc` instruction.
7. `gidata dA, dB, imm6`
- Name: "get interrupts-mode data"
 - For this instruction, the destination, mode-specific DLARs are treated as vectors of 128-bit data, where each vector element stores the following:

- * The full address (64-bit) of interrupts-mode DLARs (starting with `dB`), stored in vector element bit range `63 . . 0`.
 - * The type tag (3-bit) of interrupts-mode DLARs (starting with `dB`), stored in vector element bit range `66 . . 64`
 - Each 128-bit vector element corresponds to one interrupts-mode DLAR’s contents.
 - This instruction grabs the full address and type tag of interrupts-mode DLARs, starting with `dB`, and including the following (consecutive) `imm6` interrupts-mode DLARs, storing the 128-bit vector elements in (mode-specific) DLARs, starting with `dA` as the base DLAR. The base mode-specific DLAR is followed by consecutive DLARs if needed to store the interrupts-mode DLAR fields.
 - If the last mode-specific DLAR is not completely filled with interrupts-mode DLAR contents, its remaining 128-bit vector elements will be filled with `0x0`.
8. `ridata dA, dB, imm6`
- Name: ”restore interrupts-mode data”
 - This instruction sort of does the reverse of the `gidata` instruction, performing the equivalent of load instructions, loading into interrupts-mode DLARs based upon the metadata stored in mode-specific DLARs. As such, this instruction can read from memory, just like load instructions.
9. `giinst dA, iB, imm6`
- Name: ”get interrupts-mode instructions”
 - For this instruction, the destination, mode-specific DLARs are treated as vectors of 64-bit data, where each vector element stores the address field of a single interrupts-mode ILAR.
 - Each 64-bit vector element corresponds to one interrupts-mode ILAR’s address.
 - This instruction grabs the full address of interrupts-mode ILARs, starting with `dB` and including the following (consecutive) `imm6` interrupts-mode ILARs, storing the 64-bit vector elements in mode-specific DLARs, starting with `dA` as the base DLAR. The base mode-specific DLAR is followed by consecutive DLARs if needed to store the interrupts-mode ILAR addresses.
 - If the last mode-specific DLAR is not completely filled with interrupts-mode ILAR addresses, its remaining 64-bit vector elements will be filled with `0x0`.
10. `riinst dA, iB, imm6`
- Name: ”restore interrupts-mode instructions”
 - This instruction performs `imm6 + 1` consecutive `fetch` operations, starting with (mode-specific) `dA` as the base register that is the source

of addresses to `fetch` instructions from. User-mode `iB` is the first ILAR that instructions will be `fetch`d into.

- Basically, this instruction takes the results from a `giinst` instruction and re-`fetch`s the instructions back into interrupts-mode ILARs.

11. `gipc dA`

- Name: "get interrupts-mode pc"
- This instruction grabs the interrupts-mode program counter value, encoded into a scalar value as
 - * `concat(ilar_number, ilar_scalar_offset)`,and stores that value into DLAR `dA`.
- Note: the type tag of mode-specific DLAR `dA` is respected during this operation.
- Note: This instruction is useful for implementing interrupts-mode jump tables that are (at least partially) stored in a bunch of ILARs.

12. `ripd dA`

- Name: "restore interrupts-mode pc"
- This instruction simply does the reverse operation of the `gipc` instruction.

13. `ei`

- This instruction sets the interrupt enable register, `ie`, to `0b1`, which enables interrupts.

14. `di`

- This instruction sets the interrupt enable register, `ie`, to `0b0`, which disables interrupts.

15. `reti dA`

- Name: "return from interrupt"
- This instruction does the following:
 - * Sets `ie` to `0b1`, enabling interrupts.
 - * Switches to user-mode from interrupts-mode.
 - * Sets the user-mode `pc` to the scalar data of user-mode DLAR `dA`

16. Reserved for future expansion.

17. `addi.s fF, dA, dB, imm6`

18. `subi.s fF, dA, dB, imm6`

19. `shti.s fF, dA, dB, imm6`

20. `muli.s fF, dA, dB, imm6`

21. `divi.s fF, dA, dB, imm6`

22. `andi.s fF, dA, dB, imm6`

23. `orri.s fF, dA, dB, imm6`

24. `xori.s fF, dA, dB, imm6`
25. `shli.s fF, dA, dB, imm6`
26. `shri.s fF, dA, dB, imm6`
27. Reserved for future expansion.
28. Reserved for future expansion.
29. Reserved for future expansion.
30. Reserved for future expansion.
31. Reserved for future expansion.
32. Reserved for future expansion.

4.8 Group 7 Instructions

- Encoding: `111a aaaa aabb bbbb biii iiij jjjj jooo`
 - a: DLAR a, dA
 - b: ILAR b (iB), or DLAR b (dB)
 - i: `i_imm6`, field i 6-bit unsigned immediate
 - j: `j_imm6`, field j 6-bit unsigned immediate
 - o: opcode
- Instruction List:
- `jzo.s dA, iB, i_imm6`
 - Name: "jump if zero, scalar"
 - If dA's *scalar* data is zero, this instruction jumps to the instruction in ILAR iB at index `i_imm6`.
- `jzo.v dA, iB, i_imm6`
 - Name: "jump if zero, vector"
 - If dA's *entire* (vector) data is zero, this instruction jumps to the instruction in ILAR iB at index `i_imm6`.
- `jnz.s dA, iB, i_imm6`
 - Name: "jump if non-zero, scalar"
 - If dA's *scalar* data is non-zero, this instruction jumps to the instruction in ILAR iB at index `i_imm6`.
- `jnz.v dA, iB, i_imm6`
 - Name: "jump if non-zero, vector"
 - If dA's *entire* (vector) data is non-zero, this instruction jumps to the instruction in ILAR iB at index `i_imm6`.

- `sel.s dA, iB, i_imm6, j_imm6`
 - Name: "select, scalar"
 - If dA's scalar data is non-zero, jump to the instruction at ILAR iB, offset i_imm6. Otherwise, if dA's scalar data is zero, jump to the instruction at ILAR iB, offset j_imm6.
- `sel.v dA, iB, i_imm6, j_imm6`
 - Name: "select, vector"
 - If dA's *entire* (vector) data is non-zero, jump to the instruction at ILAR iB, offset i_imm6. Otherwise, if dA's *entire* (vector) data is zero, jump to the instruction at ILAR iB, offset j_imm6.
- Reserved for future expansion.
- Reserved for future expansion.